

Problem Set 3

Name: Gabriel Chiong

Collaborators: None

Problem 3-1.

- (a) The slots of the hash table are labelled from $0, \dots, 1$ vertically on the left-most side, with chained elements linked by an arrow.

$[0] \longrightarrow 36 \longrightarrow 92$

$[1]$

$[2]$

$[3]$

$[4] \longrightarrow 56$

$[5] \longrightarrow 47 \longrightarrow 61 \longrightarrow 33$

$[6] \longrightarrow 52$

- (b) We know the c must be at least 7, otherwise there will be at least one collision by the Pigeonhole Principle (since there are 7 elements in total). Therefore, trying integers one-by-one from $c \geq 7$ eventually yields a solution of $c = 13$.

Problem 3-2.

- (a) To hash into the same room, we choose k_1, k_2 such that $k_1 \equiv k_2 \pmod{n}$. For example, if we choose $k_1 = 3$ and $k_2 = 2n + 3$, we obtain $h(k_1) = (3a + b) \pmod{n} = h(k_2)$.
- (b) Since $u \gg n$, most adjacent k 's will cause their respective $h(k)$'s to be rounded down to the same integer. For example, for $k_1 = 1, k_2 = 2$, we have $h(k_1) = a = h(k_2)$. The $+a$ and \pmod{n} would not have a significant impact on the outcome of the hash function since $\lfloor \frac{kn}{u} \rfloor$ will always be an integer from $0, \dots, n-1$. In fact, they will have the same value regardless of the hash function chosen in \mathcal{H} .
- (c) From the lecture on hashing, the probability that two key hashes collide given a uniformly random selection of $h \in \mathcal{H}$ from a universal hash family is at most $\frac{1}{n}$, with n being the number of slots that can be hashed to. In this case, it cannot be guaranteed that they will hash to the same room.

Problem 3-3.

- (a) Each identifier can be stored in memory as a contiguous sequence of $16 \lceil \log_4(\sqrt{n}) \rceil \times 8$ bits. Therefore, we can interpret these strings as a number between $0, \dots, 2^{16 \lceil \log_4(\sqrt{n}) \rceil \times 8}$, where $2^{16 \lceil \log_4(\sqrt{n}) \rceil \times 8} = O(\sqrt{n}^{16(\log_4 2) \times 8}) = O(n^{\frac{1}{2} \cdot 16 \cdot \frac{1}{2} \cdot 8}) = O(n^{32}) = O(n^{O(1)})$. Therefore, using radix sort would allow the sorting to be done in worst-case $O(n + n \log_n n^{O(1)}) = O(n)$ time.
- (b) Effectively, each age is an integer, so we can use counting sort with a worst-case time complexity of $O(n + u) = O(n + 800,000) = O(n)$.
- (c) Multiply each integer by n^3 to obtain integers in the range of $[0, 4n^3]$. Then apply radix sort in worst-case $O(n + n \log_n n^3) = O(n)$ time.
- (d) Since this requires comparing two slices, any comparison sort is lower bounded by $\Omega(n \log n)$. Therefore, we can use merge sort for a worst-case time complexity of $\Theta(n \log n)$.

Problem 3-4.

- (a) To determine a close pair that fulfils r in expected $O(n)$ time, we first iterate through the boxes B and hash each box, b_i as the key, and its index, i as the corresponding value. Let us call the resulting hash table H . We then iterate over B a second time, each time checking H if the value $r - b_i$ exists in H . If it does exist (let us call this point b_j), check if the index value, j stored by the key $r - b_i$ in H is less than $n/10$ from i . If this holds, return our close pair. Otherwise, continue with the loop.

This algorithm makes two passes over B , each time doing expected constant work (arithmetic and hash table operations), therefore the overall worst-case time complexity is expected $O(n)$ time.

- (b) Make a first pass over all boxes and cast out boxes which have a value greater or equal to n^2 . This can be done in $O(n)$ time. Next, initialize an empty array A , and make a second pass over the new B array. At each loop iteration, append the tuple (b_i, i) to A . This can be also done in $O(n)$ time ($O(1)$ to initialize an empty array and $O(n)$ to make the second iteration).

Since each b_i is a positive integer, we can apply radix sort to array A , using the first element of the tuple, b_i as the key. We can then sort A by the number of reams in each box in $O(n + n \log_n n^2) = O(n)$ worst-case time.

Now use the "two-finger" algorithm with $i = 0$ and $j = |A| - 1$ on A . At each iteration, check the following cases:

- If $b_i + b_j = r$, check if $|i - j| < n/10$. If this holds, return true. Otherwise, continue with the algorithm by decrementing j (arbitrarily, it could have been increasing i too).
- If $b_i + b_j < r$, increase i since the boxes are ordered by their number of reams in increasing order.
- If $b_i + b_j > r$, decrease j since the boxes are ordered by their number of reams in decreasing order.

The "two-finger" algorithm touches each element at most once, therefore it runs in worst-case $O(n)$ time. If we terminate the loop with $i \geq j$, return false - we did not find a close pair that fulfilled order r . Since the overall algorithm only makes a constant number of iterations over arrays A and B (both with length at most n), the whole algorithm runs in worst-case $O(n)$ time.

Problem 3-5.

- (a) Initialize a hash table H which will map a frequency table for a particular substring in A to a count, the number of anagrams for a particular sequence of characters in A (they will have the same frequency table since they have the exact same number and type of characters). The frequency table can be implemented as a 26-tuple, an entry for every lower-case letter of the English alphabet.

To fill H , we iterate over A in two distinct parts. For the first k elements, compute their frequency table directly and insert it into H once the k -th iteration of the loop has been reached. For the next $|A| - k$ elements, use a sliding window technique to insert a new frequency table at every iteration of the loop, with the difference being removing (or decrementing) $A[i]$, and inserting (or incrementing) $A[i + k]$.

The time complexity of filling H is the sum of $O(k)$ for the first k elements, $(|A| - k) \times O(1)$ for the remaining $|A| - k$ elements (each element in the loop doing $O(1)$ work). Therefore, the overall time complexity is $O(|A|)$.

Note that since the frequency of any character is at most n , each frequency table can be thought of as a $26 \lceil \log n \rceil$ sized integer, which fits into a constant number of machine words. Therefore, it can be used as a hash key.

To check the count of B 's anagrams in $O(k)$ time, simply compute B 's frequency table in $O(k)$ time, and check if B 's frequency table exists in H . If it exists, return the corresponding count, otherwise, return 0.

- (b) Initialize result array A and the data structure from Part (a) in $O(T)$ time, with parameters T and k . Let us call this data structure H . Then iterate through \mathcal{S} . For each $S_i \in \mathcal{S}$, use $O(k)$ time to compute the frequency table of S_i (let us call this frequency table F_i), and $O(1)$ time to append the result of $H[F_i]$, the count of anagrams in T , to the result array A .

The loop runs for n iterations, therefore the time complexity for generating A is $O(nk)$. The total time complexity for the algorithm is the sum of the time to generate the data structure H and the loop over \mathcal{S} . This is $O(T + nk)$ as required.

- (c) Submit your implementation to `alg.mit.edu`.